

Causes and effects in Dichotomous Comparative Judgments: an information-theoretical system of plausible mechanism

Jose Manuel Rivera Espejo^{a,*}, Tine van Daal^a, Sven De Maeyer^a, Steven Gillis^b

^a*University of Antwerp, Training and education sciences,*

^b*University of Antwerp, Linguistics,*

Abstract

Dichotomous Comparative Judgment (DCJ) requires judges to compare pairs of stimuli to determine which one exhibits a higher degree of a specific trait. DCJ has proven effective and reliable across various fields (Pollitt, 2012b; Jones, 2015; van Daal et al., 2019; Bartholomew et al., 2018; Lesterhuis, 2018; Bartholomew and Williams, 2020; Marshall et al., 2020; Boonen et al., 2020). However, despite the method's widespread use, existing literature lacks a clear explanation of the complexities and assumptions underpinning the DCJ system, as well as the plausible mechanisms through which DCJ data could be generated. This study addresses these issues by representing DCJ within the framework of causal inference. Specifically, utilizing the structural approach, the study develops a scientific model to clarify plausible causal assumptions and mechanisms inherent in the DCJ system. The study then translates this model into a probabilistic statistical model to estimate statistical relationships and infer causal effects within the system. This research provides a robust probabilistic foundation for the statistical analysis of DCJ data, building upon Thurstone's law of comparative judgment (1927). Its findings offer valuable insights for researchers and analysts designing and implementing DCJ experiments.

Keywords: causal inference, probability, Thurstone, comparative judgement, directed acyclic graph, structural causal models, statistical modeling

1. Introduction

In contemporary contexts, Thurstone's law of comparative judgment (1927) primarily refers to the method of *dichotomous* comparative judgment (DCJ, Pollitt, 2012a,b). In DCJ, a judge assesses the relative manifestation of a *trait* within a pair of stimuli. This assessment results in a dichotomous value indicating which stimulus possesses a

*Corresponding author

Email addresses: JoseManuel.RiveraEspejo@uantwerpen.be (Jose Manuel Rivera Espejo),
tine.vandaal@uantwerpen.be (Tine van Daal), sven.demaeyer@uantwerpen.be (Sven De Maeyer),
steven.gillis@uantwerpen.be (Steven Gillis)

Preprint submitted to *Psychometrika*

August 16, 2024

higher degree of the trait. After different judges perform multiple rounds of pairwise comparisons, an outcome vector is produced. This vector is modeled using the Bradley-Terry-Luce model (BTL, [Bradley and Terry, 1952](#); [Luce, 1959](#)), which creates a score that corresponds with the trait of interest. This score is then used to rank the stimuli from lowest to highest or to evaluate the influence of certain variables on the stimuli's positions in the ranking.

DCJ has proven effective in assessing competencies and traits predominantly within the educational realm, as demonstrated by [Pollitt \(2012b\)](#), [Jones \(2015\)](#), [van Daal et al. \(2019\)](#), [Bartholomew et al. \(2018\)](#), [Lesterhuis \(2018\)](#), [Bartholomew and Williams \(2020\)](#), and [Marshall et al. \(2020\)](#). However, its application transcends education, as exemplified by [Boonen et al. \(2020\)](#). The methodology has also evolved to include multiple, as opposed to pairwise comparisons ([Luce, 1959](#); [Plackett, 1975](#)), and to accommodate comparisons with ordinal outcomes ([Tutz, 1986](#); [Agresti, 1992](#)). Overall, research suggests that DCJ offers an alternative and efficient approach to measurement and evaluation, characterized by its reliability and validity ([Lesterhuis, 2018](#); [van Daal, 2020](#); [Marshall et al., 2020](#); [Bouwer et al., 2023](#)). Nevertheless, despite the method's widespread use, existing literature lacks a clear representation of the plausible mechanisms through which DCJ data could be generated. Particularly, there is no depiction of the complexity and the assumptions underpinning the DCJ system, nor how different assessment factors can potentially influence the observed DCJ outcome.

According to [Verhavert et al. \(2019\)](#) and [van Daal \(2020\)](#), several assessment factors interact and influence the method's outcome. These factors include the number and characteristics of the stimuli, their *proximity* in terms of the assessed trait, the number of comparison per stimulus, and the pairing algorithm used. Furthermore, since the method relies on judges' assessments, the number and characteristics of judges, their *discrimination* abilities, and the number of comparisons per judge also play pivotal roles. Moreover, when the stimuli represent sub-units of higher-levels units, factors such as the number and characteristics of these units, along with their *proximity* in terms of the assessed trait, can significantly influence the outcome. For example, in [van Daal et al. \(2019\)](#), the authors assessed the academic writing skills of university students (units) using multiple argumentative essays (sub-units).

Although several studies have examined the individual impact of these factors on the method's reliability ([Bramley, 2015](#); [Pollitt, 2012b](#); [Bramley and Vitello, 2019](#); [Verhavert et al., 2019](#); [Crompvoets et al., 2022](#); [van Daal et al., 2017](#); [Gijzen et al., 2021](#); [Bouwer et al., 2023](#)), none, to the best of the authors' knowledge, have provided a transparent depiction of the DCJ system and the mechanisms generating the DCJ outcome. This study aims to fill this gap by representing DCJ within the framework of causal inference. Specifically, utilizing the structural approach to causal inference, the study develops a scientific model to clarify plausible causal assumptions and mechanisms inherent in the DCJ system. The study then translates the scientific model into a probabilistic statistical model. This model aims to produce statistical estimates to draw inferences about plausible causal relationships within the DCJ system.

Ultimately, this study provides a robust causal and probabilistic foundation for the statistical analysis of DCJ data, building upon Thurstone's law of comparative judgment ([1927](#)). Consequently, its findings offer valuable insights for researchers and analysts

designing and implementing DCJ experiments.

2. Theoretical framework

This section introduces fundamental concepts in causal inference but does not offer a comprehensive description of causal inference methods. Readers interested in deeper exploration should consult introductory papers like [Pearl \(2010\)](#), [Rohrer \(2018\)](#), [Pearl \(2019\)](#), and [Cinelli et al. \(2020\)](#). They may also find introductory books such as [Pearl and Mackenzie \(2018\)](#), [Neal \(2020\)](#) and [McElreath \(2020\)](#) useful. For more advanced study, seminal intermediate papers like [Neyman \(1923\)](#), [Rubin \(1974\)](#), [Spirtes et al. \(1991\)](#), and [Sekhon \(2009\)](#), as well as books like [Pearl \(2009\)](#), [Morgan and Winship \(2014\)](#) and [Hernán and Robins \(2020\)](#) are recommended.

2.1. The structural approach to causal inference

Empirical research addresses real-world challenges by relying on evidence gathered through observation and experimentation. In this context, researchers typically frame their research questions as *estimands* or *targets of inference*. These estimands represent the specific quantities the study aims to determine ([Everitt and Skrondal, 2010](#)). For instance, a study might examine the question, “To what extent do different teaching methods (T) influence students’ conceptual understanding of a topic (Y)?” To investigate this, the study could randomly assign students to two groups, each using a different teaching method ($T_i = \{1, 2\}$). Students’ conceptual understanding of the topic could be evaluated through pairwise comparisons, resulting in a dichotomous outcome ($Y_i = \{0, 1\}$), indicating which student among those compared has a higher level of understanding. The research question could be then framed as the estimand, “On average, is there a difference in conceptual understanding of the topic between the two groups of students?” This estimand could be mathematically expressed by the associational quantity $E[Y_i|T_i = 1] - E[Y_i|T_i = 2]$, where $E[\cdot]$ denotes the expected value. An example of this approach is seen in [Jones et al. \(2019\)](#).

Researchers then proceed to identify the estimands. *Identification* refers to the process of accurately computing an estimand using an estimator. An *estimator* is a method or function that maps data into an estimate ([Neal, 2020](#)). *Estimates* are numerical values that approximate the estimand and are derived through *estimation*, which refer to the process of integrating data with an estimator ([Everitt and Skrondal, 2010](#)). Although various methods can approximate an estimand, researchers prioritize estimators with desirable properties that ensure the accuracy of estimates. For instance, the BTL model ([Bradley and Terry, 1952](#); [Luce, 1959](#)) is an estimator known for its effectiveness in modeling pairwise comparisons, yielding accurate estimates for the target of inference discussed earlier, as long as the assumptions of the model are met.

However, many studies aim to understand the mechanisms underlying specific data types and also seek to establish causal relationships rather than merely infer associations. In the earlier example, the differences between groups obtained using the BTL model, referred to as the associational estimates, can be interpreted as causal because the data were collected through a randomized experiment. Randomized experiments enable the causal interpretation of associational estimates by ensuring several key properties, including common support, no interference, and consistency ([Morgan and Winship, 2014](#); [Neal,](#)

2020). The most crucial property, however, is that randomization effectively eliminates confounding. *Confounding* occurs when an external variable influences both the outcome and the variable of interest, leading to spurious associations (Everitt and Skrondal, 2010). Randomization mitigates this issue by decoupling the intervention assignment mechanism, such as assigning students to different groups, from other variables and outcomes (Morgan and Winship, 2014; Neal, 2020).

Experiments are widely regarded as the gold standard in evidence-based science (Hartton and Locascio, 2018; Hansson, 2014), however, researchers often encounter constraints that limit their ability to conduct experimental studies. These constraints include ethical concerns, such as the assignment of individuals to potentially harmful interventions, and practical limitations, such as the infeasibility of, for example, assigning individuals to genetic modifications or physical impairments (Neal, 2020). In these instances, causal inference offers a valuable alternative for generating causal estimates. *Causal inference* is a framework designed to identify the causes of phenomena and estimate their effects using data (Shaughnessy et al., 2010; Neal, 2020). Unlike classical statistical modeling, which focuses primarily on summarizing data and inferring associations, causal inference provides a rigorous mathematical framework for analyzing causes and counterfactuals (Pearl, 2009). Consequently, the framework offers significant theoretical insights that enhance the design of observational and experimental studies (McElreath, 2020).

While a detailed explanation of causes and counterfactuals is beyond the scope of this document, it is important to note that causal inference addresses the *fundamental problem of causality* through the use of counterfactuals (Neal, 2020). The framework defines *individual causal effects* (ICE) as the difference between an individual’s observed and unobserved potential outcomes. These unobserved potential outcomes, known as *counterfactuals*, represent hypothetical scenarios that are *contrary to fact*, where alternative outcomes for an individual, resulting from a specific cause, are neither observed nor observable (Neal, 2020; Counterfactual, 2024). The framework then extends this estimand to *average causal effects* (ACE), which captures the difference between sample averages of observed potential outcomes and counterfactuals. Finally, much like in an experiment, the framework identifies the ACE from associational estimates by conditioning for a *sufficient adjustment set* of variables, ensuring the absence of confounding (Morgan and Winship, 2014).

Several approaches to causal inference and counterfactuals exist, but two are particularly prominent: the potential outcomes approach, also known as the Neyman-Rubin causal model (Neyman, 1923; Rubin, 1974), and the structural approach (Pearl, 2009; Pearl et al., 2016). Both approaches employ rigorous mathematical notation to characterize the ACE, but they do so in different ways (Neal, 2020). The potential outcomes approach relies on counterfactual notation, whereas the structural approach employs do-calculus (Pearl, 2009). Despite these differences, both notations can be expressed in terms of the other, and both approaches provide methods for using experimental and observational data to estimate causal effects (Pearl, 2010).

The structural approach, however, offers a notable advantage over the potential outcomes approach by allowing the graphical and formal representation of theories through directed acyclic graphs (DAG, Pearl, 2009; Pearl et al., 2016; Gross et al., 2018; Neal, 2020). DAGs function as heuristics, effectively conveying the presumed causal structure

of a system, referred to as a *scientific model*. They do not represent detailed statistical models but allow researchers to deduce which statistical models can provide valid causal inferences, assuming the causal structure depicted in the DAGs are accurate (McElreath, 2020). The Identification-Estimation flowchart in Figure 1 visually represents this process of transitioning from estimands to estimates, as well as the application of the scientific model and data to identify and estimate causal effects.




Figure 1: Identification-Estimation flowchart. Extracted and slightly modified from Neal (2020, 32)

2.2. DAGs and SCMs

Graph theory is a branch of mathematics focused on the study of graphs. Graphs are mathematical structures modeling pairwise relations between objects. They can represent physical relations, such as electrical circuits and roadways, and less tangible structures, such as ecosystems and sociological relations. Graphs have proven useful in various fields, including computer science, operations research, and the natural and social sciences (Gross et al., 2018).

In statistics, one application incorporating concepts from graph theory is causal inference. Specifically, the structural approach to causal inference uses directed acyclic graphs (DAG) to provide a graphical and formal representation of the causal structure of a system (Neal, 2020). In this context, a *graph* denotes a collection of nodes connected by edges, where nodes represent random variables. The term *directed* indicates the edges of the graph extend from one node to another, with arrows showing the direction of causal influence. Moreover, the term *acyclic* indicates the causal influences do not form a loop, meaning the influences do not cycle back on themselves (McElreath, 2020).

DAGs offer two key advantages for modeling causal structures. Firstly, they represent causal relations in a nonparametric and fully interactive manner. This feature allows for feasible causal analysis strategies without needing the specification of the type of data or the nature of the functional dependence among variables (Morgan and Winship, 2014). Secondly, regardless of complexity, DAGs can represent various causal structures using only five fundamental building blocks (Neal, 2020; McElreath, 2020). This feature enables the decomposition of complex structures into basic building blocks, facilitating the analysis of these structures by focusing on the causal assumptions associated with each building block individually (McElreath, 2020). These building blocks can be repre-

sented in three ways: the magnified representation, the standard representation, and the structural causal model form (SCM, Morgan and Winship, 2014).

The left panels of Figure 2 illustrate the *magnified* representation. These graphs depict the *endogenous* variables $V = \{X, Z, Y\}$ alongside the *exogenous* variables $E = \{e_X, e_Z, e_Y\}$. Endogenous variables are those whose causal mechanisms the investigator chooses to model (Neal, 2020). In contrast, exogenous variables represent *errors* or *disturbances* arising from omitted factors that the investigator chooses not to model explicitly (Pearl, 2009, 27,68). The graphs show endogenous variables as solid black circles to signify that they are observed random variables, while endogenous variables are depicted as open circles to signify their unobserved (latent) nature. Lastly, the arrows in the graphs reflect the expected direction of causal influences among these variables.

Often, DAGs omit the exogenous variables for simplicity, resulting in the *standard* representation. However, including exogenous variables in a graph can be beneficial in some scenarios, as their presence can reveal potential issues related to conditioning and confounding (Cinelli et al., 2020), concepts explored in Section 2.3. The standard representation is illustrated in the middle panels of Figure 2.

Lastly, the right panels of Figure 2 depict the SCM form of the fundamental building blocks. SCMs are formal mathematical models defined by a set of endogenous variables V , a set of exogenous variables E , and a set of functions $F = \{f_X, f_Z, f_Y\}$ (Pearl, 2009; Neal, 2020). These functions, referred to as structural equations, specify each endogenous variable as nonparametric functions of other variables. Moreover, SCMs use the symbol ‘:=’ to indicate the variables’ asymmetrical causal dependence and the symbol ‘ $\perp\!\!\!\perp$ ’ to represent *d-separation*, which roughly equates to the concept of variable independence. The concepts of d-separation and causal (in)dependence are explored in Section 2.3.

A careful examination of Figure 2 highlights the assumptions underlying these building blocks. Figures 2a, 2b, and SCM 2c depict two unconnected nodes, representing a scenario where variables X and Y are not causally related. Figures 2d, 2e, and SCM 2f illustrate two connected nodes, showing a scenario where a *parent* node X exerts a causal influence on a *child* node Y . Consequently, Y is considered a *descendant* of X . Figures 2g, 2h, and SCM 2i depict a *chain* or *pipe*, where X influences Z , and Z influences Y . In this configuration, X is a parent node of Z , and Z is a parent node of Y . This creates a *directed path* between X and Y . Consequently, X is an *ancestor* of Y , and Z fully *mediates* the relationship between the two. Figures 2j, 2k, and SCM 2l illustrate a *fork*, where variables X and Y are both influenced by Z . Here, Z is a parent node of X and Y . Finally, Figures 2m, 2n, SCM 2o depict a *collider*, also known as *inmortality*, where variables X and Y are concurrent causes of Z . In this configuration, X and Y are not causally related to each other but both influence Z . Additionally, in all SCMs, the errors are assumed to be mutually independent of each other and of all other variables in the graph, as evidenced by the pairwise relations $e_X \perp\!\!\!\perp e_Y$, $e_X \perp\!\!\!\perp e_Z$, and $e_Z \perp\!\!\!\perp e_Y$.

The motivating example in Section 2.1 can further illustrate how to use the five fundamental building blocks to construct a system’s causal structure. In this scenario, the investigator aims to determine whether, *on average*, there is a difference in conceptual understanding of a topic ($Y_i = \{0, 1\}$) between two groups of students ($T_i = \{1, 2\}$), described by the estimand $E[Y_i|T_i = 1] - E[Y_i|T_i = 2]$. However, unlike the previous

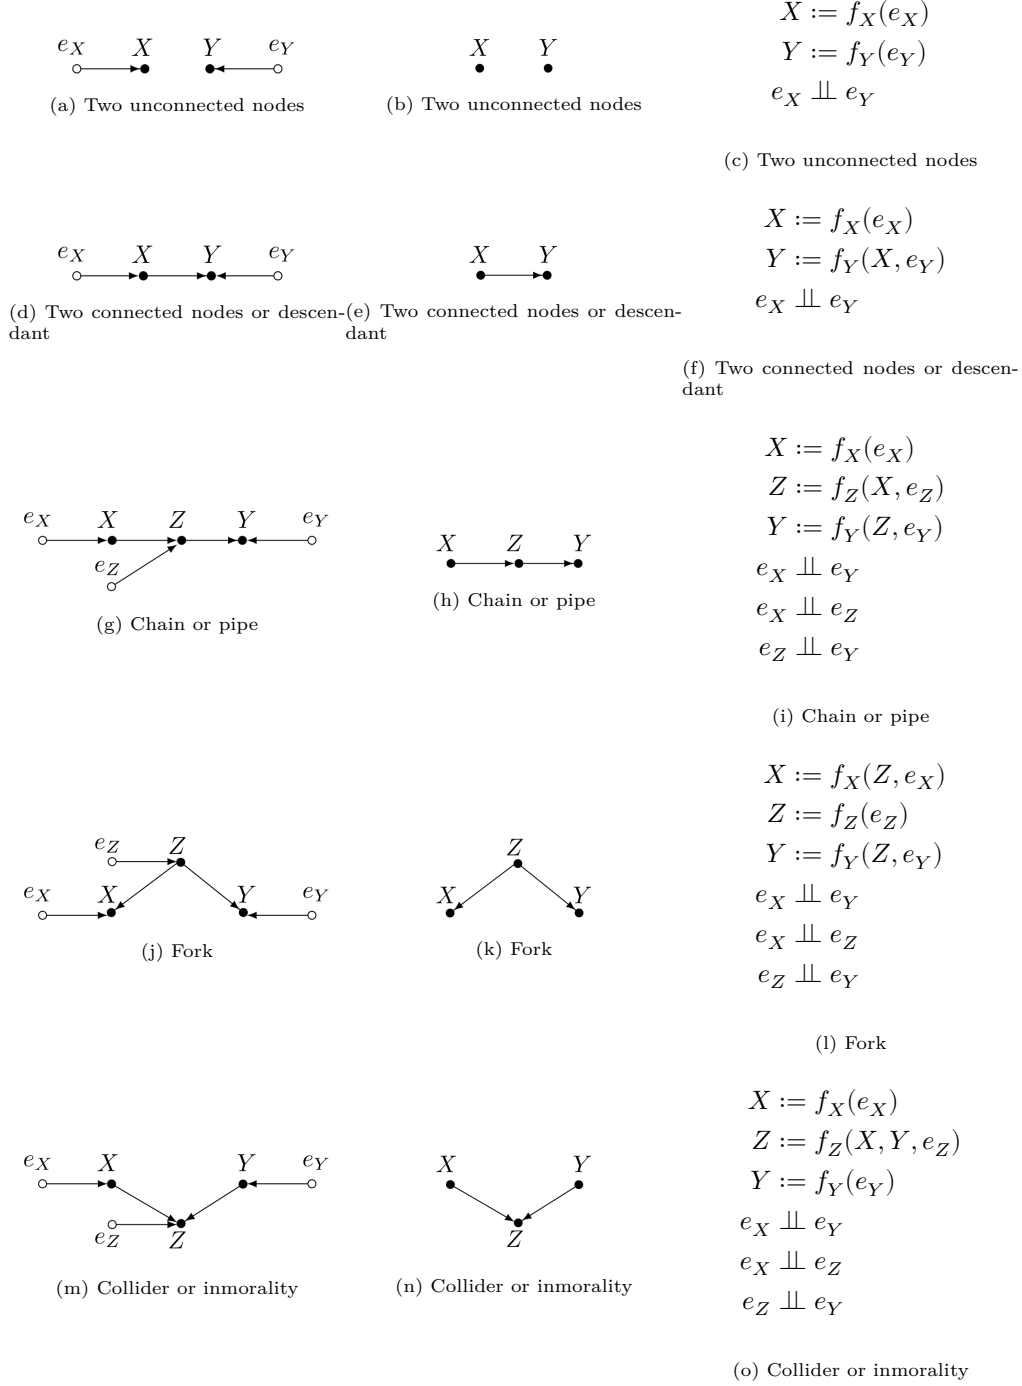


Figure 2: The five fundamental building blocks of DAGs. **Note:** left panels show the magnified representation, middle panels show the standard representation, and the right panels show their corresponding SCM form.

case, an experiment cannot be conducted, and the problem suggests that the country to which a student belongs (X) may influence both T and Y . Such scenarios are plausible, especially when the teaching methods depend on software or access to technology, which may be limited in certain countries (maybe an example with more impact?). Figure 3 illustrates the plausible causal structure of the motivating example. A detailed examination of Figures 3a, 3b, and SCM 3c reveals the presence of at least four of the five fundamental building blocks. The figures display multiple descendants, as indicated by pairwise relations such as $X \rightarrow T$, $X \rightarrow Y$, and $T \rightarrow Y$. Additionally, the figures features multiple pairs of unconnected nodes, evident from the relations $e_T \perp\!\!\!\perp e_X$, $e_T \perp\!\!\!\perp e_Y$, and $e_X \perp\!\!\!\perp e_Y$. Finally, the figures illustrate the fork $X \rightarrow \{T, Y\}$, and two colliders with $\{X, e_T\} \rightarrow T$ and $\{X, T, e_Y\} \rightarrow Y$.




Figure 3: DAGs for a plausible causal structure in a system.

2.3. The flow of association and causation in graphs




Figure 4: The flow of association and causation in graphs. Extracted and slightly modified from Neal (2020, 31)

3. Theory

3.1. A scientific model for the DCJ

3.2. Probabilistics assumptions of the scientific model

3.3. From the scientific to statistical model

3.4. Let's talk about Thurstone

4. Discussion

4.1. Findings

4.2. Limitations and further research

5. Conclusion

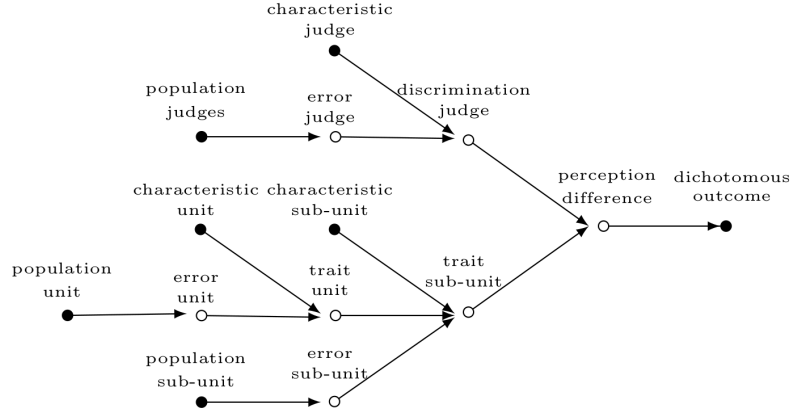


Figure 5: DCJ causal diagram, simplified description




Figure 6: DCJ causal diagram, simplified mathematical description




Figure 7: DCJ causal diagram, population mathematical description




Figure 8: DCJ causal diagram, sample with comparisons mathematical description

Declarations

Funding: The project was founded through the Research Fund of the University of Antwerp (BOF).

Financial interests: The authors have no relevant financial interest to disclose.

Non-financial interests: Author XX serve on advisory board of Company Y but receives no compensation this role.

Ethics approval: The University of Antwerp Research Ethics Committee has confirmed that no ethical approval is required.

Consent to participate: Not applicable

Consent for publication: All authors have read and agreed to the published version of the manuscript.

Availability of data and materials: No data was utilized in this study.

Code availability: All the code utilized in this research is available in the digital document located at: https://jriverspejo.github.io/paper2_manuscript/.

Authors' contributions: *Conceptualization:* S.G., S.D.M., T.vD., and J.M.R.E.; *Methodology:* S.D.M., T.vD., and J.M.R.E.; *Software:* J.M.R.E.; *Validation:* J.M.R.E.; *Formal Analysis:* J.M.R.E.; *Investigation:* J.M.R.E.; *Resources:* S.G., S.D.M., and T.vD.; *Data curation:* J.M.R.E.; *Writing - original draft:* J.M.R.E.; *Writing - review & editing:* S.G., S.D.M., and T.vD.; *Visualization:* J.M.R.E.; *Supervision:* S.G. and S.D.M.; *Project administration:* S.G. and S.D.M.; *Funding acquisition:* S.G. and S.D.M.

6. Appendix

- 6.1. Why do we need to estimate judges' abilities?*
- 6.2. Latent variables as a mean of imputation*
- 6.3. Other comparative scenarios*

References

- Agresti, A., 1992. Analysis of ordinal paired comparison data. *Journal of the Royal Statistical Society* 41, 287–297. URL: <https://www.jstor.org/stable/2347562>, doi:10.2307/2347562.
- Bartholomew, S., Nadelson, L., Goodridge, W., Reeve, E., 2018. Adaptive comparative judgment as a tool for assessing open-ended design problems and model eliciting activities. *Educational Assessment* 23, 85–101. doi:10.1080/10627197.2018.1444986.
- Bartholomew, S., Williams, P., 2020. Stem skill assessment: An application of adaptive comparative judgment, in: Anderson, J., Li, Y. (Eds.), *Integrated Approaches to STEM Education*. Advances in STEM Education. Springer, pp. 331–349. doi:10.1007/978-3-030-52229-2_18.
- Boonen, N., Kloots, H., Gillis, S., 2020. Rating the overall speech quality of hearing-impaired children by means of comparative judgements. *Journal of Communication Disorders* 83, 1675–1687. doi:10.1016/j.jcomdis.2019.105969.
- Bouwer, R., Lesterhuis, M., De Smedt, F., Van Keer, H., De Maeyer, S., 2023. Comparative approaches to the assessment of writing: Reliability and validity of benchmark rating and comparative judgement. *Journal of Writing Research* 15, 497–518. URL: <https://www.jowr.org/index.php/jowr/article/view/867>, doi:10.17239/jowr-2024.15.03.03.
- Bradley, R., Terry, M., 1952. Rank analysis of incomplete block designs: I. the method of paired comparisons. *Biometrika* 39, 324–345. URL: <http://www.jstor.com/stable/2334029>, doi:10.2307/2334029.
- Bramley, T., 2015. Investigating the reliability of adaptive comparative judgment. URL: <http://www.cambridgeassessment.org.uk/Images/232694-investigating-the-reliability-of-adaptive-comparative-judgment.pdf>. cambridge Assessment Research Report.
- Bramley, T., Vitello, S., 2019. The effect of adaptivity on the reliability coefficient in adaptive comparative judgement. *Assessment in Education: Principles, Policy and Practice* 71, 1–25. doi:10.1080/0969594X.2017.1418734.
- Cinelli, C., Forney, A., Pearl, J., 2020. A crash course in good and bad controls. SSRN URL: <https://ssrn.com/abstract=3689437>, doi:10.2139/ssrn.3689437.
- Counterfactual, 2024. Merriam-webster.com dictionary. URL: <https://www.merriam-webster.com/dictionary/hacker>. retrieved July 23, 2024.
- Crompvoets, E.A.V., Béguin, A.A., Sijtsma, K., 2022. On the bias and stability of the results of comparative judgment. *Frontiers in Education* 6. URL: [url{https://www.frontiersin.org/articles/10.3389/feduc.2021.788202}](https://www.frontiersin.org/articles/10.3389/feduc.2021.788202), doi:10.3389/feduc.2021.788202.
- Everitt, B., Skrondal, A., 2010. *The Cambridge Dictionary of Statistics*. Cambridge University Press.
- Gijzen, M., van Daal, T., Lesterhuis, M., Gijbels, D., De Maeyer, S., 2021. The complexity of comparative judgments in assessing argumentative writing: An eye tracking study. *Frontiers in Education* 5. URL: [url{https://www.frontiersin.org/articles/10.3389/feduc.2020.582800}](https://www.frontiersin.org/articles/10.3389/feduc.2020.582800), doi:10.3389/feduc.2020.582800.
- Gross, J., Yellen, J., Anderson, M., 2018. *Graph Theory and Its Applications*. Textbooks in Mathematics, Chapman and Hall/CRC. doi:<https://doi.org/10.1201/9780429425134>. 3rd edition.
- Hansson, S., 2014. Why and for what are clinical trials the gold standard? *Scandinavian Journal of Public Health* 42, 41–48. doi:10.1177/1403494813516712. PMID: 24553853.
- Hariton, E., Locascio, J., 2018. Randomised controlled trials – the gold standard for effectiveness research. *BJOG: An International Journal of Obstetrics & Gynaecology* 125, 1716–1716. URL: <https://obgyn.onlinelibrary.wiley.com/doi/abs/10.1111/1471-0528.15199>, doi:10.1111/1471-0528.15199.
- Hernán, M., Robins, J., 2020. *Causal Inference: What If*. 1 ed., Chapman and Hall/CRC. URL: <https://www.hsph.harvard.edu/miguel-hernan/causal-inference-book>. last accessed 31 July 2024.
- Jones, I., 2015. The problem of assessing problem solving: can comparative judgement help? *Educational Studies in Mathematics* 89, 337–355. doi:10.1007/s10649-015-9607-1.
- Jones, I., Bisson, M., Gilmore, C., Inglis, M., 2019. Measuring conceptual understanding in randomised controlled trials: Can comparative judgement help? *British Educational Research Journal* 45, 662–680. URL: <https://bera-journals.onlinelibrary.wiley.com/doi/abs/10.1002/berj.3519>, doi:10.1002/berj.3519.
- Lesterhuis, M., 2018. The validity of comparative judgement for assessing text quality: An assessor’s perspective. Ph.D. thesis. University of Antwerp.
- Luce, R., 1959. On the possible psychophysical laws. *The Psychological Review* 66, 482–499. doi:10.1037/h0043178.
- Marshall, N., Shaw, K., Hunter, J., Jones, I., 2020. Assessment by comparative judgement: An application to secondary statistics and english in new zealand. *New Zealand Journal of Educational Studies* 55, 49–71. doi:10.1007/s40841-020-00163-3.

- McElreath, R., 2020. *Statistical Rethinking: A Bayesian Course with Examples in R and STAN*. Chapman and Hall/CRC.
- Morgan, S., Winship, C., 2014. *Counterfactuals and Causal Inference: Methods and Principles for Social Research*. Analytical Methods for Social Research. 2 ed., Cambridge University Press.
- Neal, B., 2020. Introduction to causal inference from a machine learning perspective. URL: https://www.bradyn Neal.com/Introduction_to_Causal_Inference-Dec17_2020-Neal.pdf. last accessed 30 April 2024.
- Neyman, J., 1923. On the application of probability theory to agricultural experiments. essay on principles. section 9. *Statistical Science* 5, 465–472. URL: <http://www.jstor.org/stable/2245382>. translated by Dabrowska, D. and Speed, T. (1990).
- Pearl, J., 2009. *Causality: Models, Reasoning and Inference*. Cambridge University Press.
- Pearl, J., 2010. An introduction to causal inference. *The international journal of biostatistics* 6, 855–859. URL: <https://www.degruyter.com/document/doi/10.2202/1557-4679.1203/html>, doi:10.2202/1557-4679.1203.
- Pearl, J., 2019. The seven tools of causal inference, with reflections on machine learning. *Communications of the ACM* 62, 54–60. doi:10.1177/0962280215586010.
- Pearl, J., Glymour, M., Jewell, N., 2016. *Causal Inference in Statistics: A Primer*. John Wiley & Sons, Inc.
- Pearl, J., Mackenzie, D., 2018. *The Book of Why: The New Science of Cause and Effect*. 1st ed., Basic Books, Inc.
- Plackett, R., 1975. The analysis of permutations. *Journal of the Royal Statistical Society* 24, 193–202. URL: <https://www.jstor.org/stable/2346567>, doi:10.2307/2346567.
- Pollitt, A., 2012a. Comparative judgement for assessment. *International Journal of Technology and Design Education* 22, 157–170. doi:10.1007/s10798-011-9189-x.
- Pollitt, A., 2012b. The method of adaptive comparative judgement. *Assessment in Education: Principles, Policy and Practice* 19, 281–300. doi:10.1080/0969594X.2012.665354.
- Rohrer, J., 2018. Thinking clearly about correlations and causation: Graphical causal models for observational data. *Advances in Methods and Practices in Psychological Science* 1, 27–42. doi:10.1177/2515245917745629.
- Rubin, D., 1974. Estimating causal effects of treatments in randomized and nonrandomized studies. *Journal of Educational Psychology* 66, 688–701. doi:10.1037/h0037350.
- Sekhon, J., 2009. The neyman-rubin model of causal inference and estimation via matching methods, in: Box-Steffensmeier, J., Brady, H., Collier, D. (Eds.), *The Oxford Handbook of Political Methodology*. Oxford University Press, pp. 271–299. doi:10.1093/oxfordhb/9780199286546.003.0011.
- Shaughnessy, J., Zechmeister, E., Zechmeister, J., 2010. *Research Methods in Psychology*. McGraw-Hill. URL: https://web.archive.org/web/20141015135541/http://www.mhhe.com/socscience/psychology/shaugh/ch01_concepts.html. retrieved July 23, 2024.
- Spirtes, P., Glymour, C., Scheines, R., 1991. From probability to causality. *Philosophical Studies* 64, 1–36. URL: <https://www.jstor.org/stable/4320244>.
- Thurstone, L., 1927. A law of comparative judgment. *Psychological Review* 34, 482–499. doi:10.1037/h0070288.
- Tutz, G., 1986. Bradley-terry-luce model with an ordered response. *Journal of Mathematical Psychology* 30, 306–316. doi:10.1016/0022-2496(86)90034-9.
- van Daal, T., 2020. Making a choice is not easy?!: Unravelling the task difficulty of comparative judgement to assess student work. Ph.D. thesis. University of Antwerp.
- van Daal, T., Lesterhuis, M., Coertjens, L., Donche, V., De Maeyer, S., 2019. Validity of comparative judgement to assess academic writing: examining implications of its holistic character and building on a shared consensus. *Assessment in Education: Principles, Policy & Practice* 26, 59–74. doi:10.1080/0969594X.2016.1253542.
- van Daal, T., Lesterhuis, M., Coertjens, L., van de Kamp, M., Donche, V., De Maeyer, S., 2017. The complexity of assessing student work using comparative judgment: The moderating role of decision accuracy. *Frontiers in Education* 2. URL: <https://www.frontiersin.org/articles/10.3389/feduc.2017.00044>, doi:10.3389/feduc.2017.00044.
- Verhavert, S., Bouwer, R., Donche, V., De Maeyer, S., 2019. A meta-analysis on the reliability of comparative judgement. *Assessment in Education: Principles, Policy and Practice* 26, 541–562. doi:10.1080/0969594X.2019.1602027.